



DDL. Data Definition Language

Unit 9



Introduction

- The DDL (Data Definition Language), is part of the SQL. It covers the statements orientated to define the database structure.
- That is why it is usually used by the database administrator, and sometimes by programmers.
- Its verbs are:
 - CREATE
 - DROP
 - ALTER
- The objects involved in the statements that we will learn are:
 - Databases
 - Tables
 - Indexes
 - Views



DDL Database

A database is stored in several files:

- The main data file .mdf
- Secondary data files .ndf
- Log files .ldf



CREATE DATABASE

```
CREATE DATABASE DBname  
[ ON [ PRIMARY ] [ <file_desc> [ ,...n ]  
[ LOG ON { < file_desc > [ ,...n ] } ] ]  
[ COLLATE CollationName]  
[ WITH <external_access_option> ] [;]
```



CREATE DATABASE

```
<file_desc> ::=  
  (NAME = LogicalFileName ,  
    FILENAME = 'OSFileName'  
    [ , SIZE = size [ KB | MB | GB | TB ] ]  
    [ , MAXSIZE = { max_size [ KB | MB | GB | TB ]  
                  | UNLIMITED } ]  
    [ , FILEGROWTH = growth_inc  
                  [ KB | MB | GB | TB | % ] ]  
  )
```



CREATE DATABASE

COLLATE <CollationName>

< CollationName >:: =

CollationDesignator_ CaseSensitivity_AccentSensitivity

- CollationDesignator : Windows collation name, specifies the base collation rules.
- CaseSensitivity: CS specifies Case-Sensitive, CI specifies Case-Insensitive
- AccentSensitivity: AS specifies Accent-Sensitive, AI specifies Accent-Insensitive

Example:

COLLATE Modern_Spanish_CS_AI



DROP DATABASE

Eliminate a database (data and structure).

```
DROP DATABASE { DBname } [ ,...n ] [;]
```



ALTER DATABASE

Modify the properties of the database.

```
ALTER DATABASE DBname
{
    <FileChanges>
    | <GroupChanges>
    | <options>
    | MODIFY NAME = new DBname
    | COLLATE collationName
}
[;]
```




ALTER DATABASE

<FileChanges> ::=

```
{ ADD FILE < file_desc > [ ,...n ]  
  [ TO FILEGROUP { groupName | DEFAULT } ]  
  | ADD LOG FILE < file_desc > [ ,...n ]  
  | REMOVE FILE FileName  
  | MODIFY FILE <file_desc>
```

<file_desc> ::=

```
( NAME = FileName [ , NEWNAME = new_FileName ]  
  [ , FILENAME = 'OS_FileName' ]  
  [ , SIZE = size [ KB | MB | GB | TB ] ]  
  [ , MAXSIZE = { max_size [ K | MB | GB | TB ] | UNLIMITED } ]  
  [ , FILEGROWTH = increment [ KB | MB | GB | TB | % ] ]  
  [ , OFFLINE ] )
```



CREATE TABLE

```
CREATE TABLE [ DBName.][SchemaName.] TableName  
  ( { <Column_def>  
    | < ComputedColumn_def> } [ ,...n ]  
    [ <table_constraint> ] [ ,...n ] ) [ ; ]
```

```
< Column_def > ::=  
  ColName <Datatype> [ NULL | NOT NULL ]  
  [ COLLATE CollationName ]  
  [ [ CONSTRAINT Constraint_name ]  
    DEFAULT constant_expr ]  
  [ IDENTITY [( seed, increment )] ]  
  [ ROWGUIDCOL ]  
  [ <Col_constraint> [ ...n ] ]
```



CREATE TABLE

```
< Col_constraint > ::= [ CONSTRAINT Constraint_name ]  
  { { PRIMARY KEY | UNIQUE } [ CLUSTERED | NONCLUSTERED ]  
    | [ FOREIGN KEY ] REFERENCES  
      [ SchemaName. ] referenced_tbl_name [ (  
  ref_col) ]  
    [ ON DELETE { NO ACTION | CASCADE | SET NULL |  
                  SET DEFAULT } ]  
    [ ON UPDATE { NO ACTION | CASCADE | SET NULL |  
                  SET DEFAULT } ] [ NOT FOR REPLICATION ]  
    | CHECK (logical_expression)  
  }
```



CREATE TABLE

```
< table_constraint > ::=  
[ CONSTRAINT Constraint_name ] {  
    {PRIMARY KEY|UNIQUE } [CLUSTERED|NONCLUSTERED ]  
        (ColName[ASC|DESC][ ,...n ] )  
    | FOREIGN KEY(ColName[ ,...n ])REFERENCES  
        referenced_tbl_name  
            [(ref_col)[ ,...n ] )]  
    [ ON DELETE { NO ACTION | CASCADE | SET NULL |  
        SET DEFAULT } ]  
    [ ON UPDATE { NO ACTION | CASCADE | SET NULL |  
        SET DEFAULT } ]  
    | CHECK (logical_expression ) }
```



CREATE TABLE

Example:

```
CREATE TABLE persons (  
    id int,  
    idcard int,  
    mother int,  
    father int,  
    tutor varchar(10),  
    PRIMARY KEY (id),  
    UNIQUE (idcard),  
    FOREIGN KEY (mother) REFERENCES Persons,  
    FOREIGN KEY (father) REFERENCES Persons (id),  
    FOREIGN KEY (tutor) REFERENCES Persons (idcard));
```



CREATE TABLE

Example:

```
CREATE TABLE LineOrder (  
    order int,  
    line int,  
    productID int,  
    supplierID int,  
    quantity int,  
    PRIMARY KEY (order, line),  
    FOREIGN KEY (productID ,supplierID) REFERENCES Products,  
    UNIQUE (order, productID ,supplierID));
```



CREATE TABLE

Computed columns:

< Column_def > ::=

ColName AS expression [PERSISTED
[NOT NULL]
[<Col_constraint> [... *n*]]]

- A computed column can be used in the select list, in the WHERE clause , and in the ORDER BY clause.
- It cannot be used in a column constraint with the DEFAULT, FOREIGN KEY, NOT NULL clause.
- It cannot be the target of an INSERT or UPDATE.



CREATE TABLE

TEMPORARY TABLES

- It is a table created in a process and when the process finishes the table is dropped.
- It can be local or global.
- Local temporary tables are available only in the current session, global temporary tables are available in all the sessions.

#TableName for a local temporary table

##TableName for a global temporary table

- Example:

```
CREATE TABLE #aux (col1 INT PRIMARY KEY);
```




DROP TABLE

- Erase a table (data and structure)

`DROP TABLE [[DBName].SchemaName.]TableName[,...n] [;]`

- Referential integrity is checked.
- Several tables can be dropped at the same time.
- Tables are dropped in the same order that they appear in the sentence. Useful when dropping referenced tables.
- Example:
`DROP TABLE orders, products`



ALTER TABLE

- Modify its definition

```
ALTER TABLE [DBName.[SchName].| SchName.]TableName
{ ALTER COLUMN ColName <columnChanges>
| [WITH{ CHECK | NOCHECK}] ADD
    { <Column_def>
      | <ComputedCol_Def>
      | <table_constraint> } [ ,...n ]
| DROP {[CONSTRAINT] constraintName
      | COLUMN ColName} [,...n ]
| {CHECK|NOCHECK} CONSTRAINT
                                   {ALL|constraint_name[,...n ]}
| {ENABLE|DISABLE} TRIGGER {ALL | TriggerNa [ ,...n ] }
}[ ; ]
```



ALTER TABLE

```
< columnChanges >::=  
{  
    <Datatype>  
        [ NULL | NOT NULL ]  
        [ COLLATE CollationName  
        | {ADD | DROP } { ROWGUIDCOL |PERSISTED }  
}
```

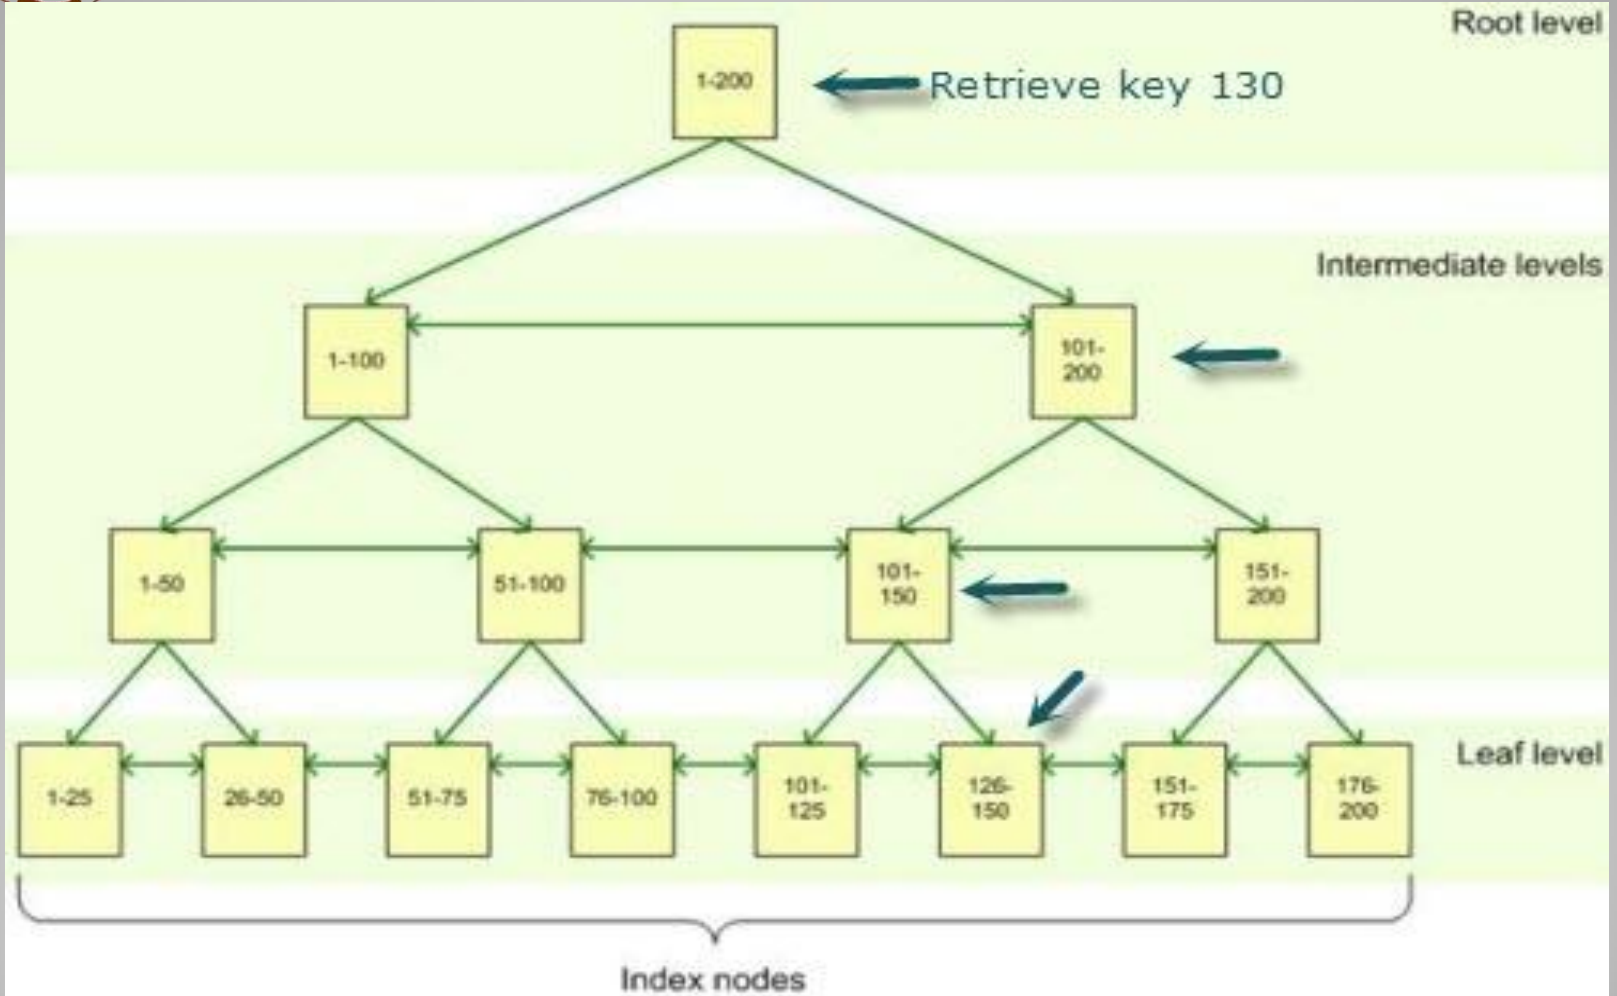


INDEXES

- An index is an on-disk structure associated with a table or view that speeds up the retrieval of rows.
- An index contains keys built from one or more columns in the table. These keys are stored in a structure (B-tree) that enables SQL Server to find the row or rows associated with the key values quickly and efficiently.
- The selection of the right indexes for a database and its workload is a complex balancing act between query speed and update cost, usually done by the DB Administrator.



INDEXES





INDEXES

- Instead of walking through *the whole table* and checking every row to see if it matches, the search is performed in the B-tree loaded in the memory.
- The leaf node contains either the entire row of data or a pointer to that row, depending on whether the index is clustered or nonclustered.
- A table that has a clustered index is referred to as a *clustered table (tabla agrupada)*. A table that has no clustered index is referred to as a *heap (montón)*.



INDEXES

INDEX TYPES

- Single column and composite index
 - Single column index → based on one column.
 - Composite index → based on several columns.
- Clustered and nonclustered index.
 - Clustered index (índice agrupado) → A clustered index stores the actual data rows at the leaf level of the index. So, the data rows are stored in order based on the index key.
 - Unclustered index → At the leaf level we have a pointer to the data row. The data rows are not guaranteed to be in any particular order.
- Unique index
 - An index that ensures the uniqueness of each value in the indexed column.



INDEXES

PROS

- Indexes can enhance performance because they can provide a quick way for the query engine to find data.
- Indexes speed up the ordering process.

CONS

- They can take up significant disk space.
- Indexes are automatically updated when the data rows themselves are updated, which can lead to additional overhead and can affect performance.



INDEXES

```
CREATE [ UNIQUE ] [ CLUSTERED | NONCLUSTERED ] INDEX  
    index_name  
    ON <object> (column_na [ ASC | DESC ] [ ,...n ] )  
[ ; ]
```

<object> ::=

```
{  
[DBName].[SchemaName].| SchemaName.]TableName  
}
```



INDEXES

Examples:

```
CREATE INDEX Iclients_name ON Clients (name)
```

```
CREATE UNIQUE INDEX Iproducts_des ON products  
(description)
```

```
CREATE UNIQUE CLUSTERED INDEX Iorders_pk ON  
orders (id)
```

```
CREATE INDEX Ioffices_sal ON Offices (sales DESC)
```

```
CREATE INDEX Ioffices _2 ON Offices (region,city)
```



INDEXES

```
DROP INDEX <index> [ ,...n ] [ ; ]
```

```
<indice> ::=
```

```
{  
indexname ON  
  [DBName.[SchemaName].|SchemaName.]TableName  
}
```

Example: DROP INDEX Iclients_name ON Clients



INDEXES

- ALTER INDEX

Modifies an existing index by:

- Disabling,
- Rebuilding,
- Reorganising (leaf level),
- setting options on it.



VIEWS

```
CREATE VIEW [SchemaName.] ViewName  
    [ (column_a [ ,...n ] ) ]  
    AS select_sentence [ ; ]
```

select_sentence:

- can use several tables and views.
- Can use functions, UNION or UNION ALL.
- Maximum 1.024 columns
- Is used as a table, but a virtual table. Its definition and execution plan are stored, which enhances performance.



VIEWS

`DROP VIEW [SchemaName.]ViewName[,...n] [;]`

- If we drop a table, before hand we will have to drop all views based on the table, because they will not be eliminated automatically.